

Sihab Safin

sihabulislamsafin1@gmail.com | github.com/sihabsafin | linkedin.com/in/sihabsafin | sihabsafin.vercel.app

SKILLS

Languages: Python, SQL, JavaScript, HTML/CSS

GenAI / LLM — Core: LLM Orchestration, Prompt Engineering, Prompt Optimization, Context Engineering, Chain-of-Thought Prompting, Model Evaluation, Hallucination Reduction, Guardrails & Safety Filtering

GenAI / LLM — Retrieval: RAG (Retrieval-Augmented Generation), Semantic Search, Hybrid Search (BM25 + Vector + RRF), Vector Embeddings, Reranking, Chunking Strategies, Query Rewriting

GenAI / LLM — Fine-Tuning: QLoRA, PEFT, LoRA, 4-bit Quantization (NF4), Instruction Tuning, Supervised Fine-Tuning (SFT), Transfer Learning

GenAI / LLM — Evaluation & Cost: RAGAS, LLM-as-Judge, Prompt Regression Testing, A/B Testing, LLM Cost Optimization, Latency Optimization

Agentic AI — Systems: Multi-Agent Systems, Agent Orchestration, Autonomous Workflows, Task Planning & Decomposition, Agent-to-Agent Communication, Shared State Management

Agentic AI — Tooling: CrewAI, Tool-Use / Function Calling, Structured Output Parsing, Workflow Automation, n8n, Webhook Integrations, Retry / Fallback Logic

LLM APIs & Models: OpenAI GPT-4o, Google Gemini, Groq, Anthropic Claude, Qwen2.5, Gemma-2, LLaMA-3.2

ML / Data Science — Modeling: Scikit-learn, Deep Learning, NLP, Model Training & Evaluation, Classification, Regression, Statistical Analysis

ML / Data Science — Data: Pandas, NumPy, Data Pipelines, ETL, Data Cleaning & Preprocessing, Excel (Dashboards, Pivot Tables)

Frameworks & Infra — AI/ML: LangChain, Hugging Face Transformers, PyTorch, Google Colab (GPU Training)

Frameworks & Infra — Backend/Data: FastAPI, Supabase, Vector Databases (Pinecone, Weaviate, FAISS), Redis, Streamlit, Docker

Frameworks & Infra — Tooling: Git/GitHub, REST API Design, CI-Friendly Development Practices, Cloud Deployment

Voice / Real-Time AI: LiveKit, Deepgram, ElevenLabs, Real-Time Voice Agents, Speech-to-Speech Pipelines

Security & Standards: MITRE ATT&CK, NIST NVD API, Threat Detection Pipelines, Vulnerability Triage

Client & Delivery: Requirements Discovery & Scoping, Stakeholder Communication, Technical Documentation for Non-Technical Audiences, End-to-End Project Delivery, Client-Facing Consulting

Other: Technical Writing / Research Documentation, SEO, WordPress/WooCommerce

EXPERIENCE

Agentic AI Engineer — Wizard Labs (wizardlabs.com)

Jul 2025 – May 2026

- Designed and shipped multi-agent AI systems using CrewAI, orchestrating specialized agents for autonomous task planning, tool-use, and multi-step workflow execution across production use cases.
- Built a 13-agent cybersecurity threat detection pipeline (AutoSEC) integrating MITRE ATT&CK mappings and the NIST NVD API across 5 sequential analysis phases, improving structured vulnerability triage.
- Developed a real-time conversational voice agent (InterviewIQ) combining LiveKit, Deepgram, ElevenLabs, and GPT-4o for low-latency speech-to-speech interaction.
- Architected a hybrid retrieval system (DocMind Hybrid) combining BM25 keyword search, dense semantic embeddings, and Reciprocal Rank Fusion (RRF) to improve document retrieval accuracy.
- Built a multi-agent investment research dashboard (InvestIQ AI), coordinating agents for data gathering, analysis, and report generation with a Streamlit front end.

Generative AI Engineer — Kernelor (kernelor.com)

Mar 2024 – Jun 2025

- Built RecruitIQ, a multi-agent hiring pipeline using CrewAI, Groq, and Gemini with Supabase for data persistence and Streamlit for the user interface, automating candidate screening workflows.
- Developed ClinicalCrew, a 4-phase clinical decision support system coordinating 11 specialized agents to assist structured medical reasoning workflows.
- Implemented retrieval-augmented generation (RAG) pipelines integrating LLM APIs (OpenAI, Gemini, Groq) with vector search for grounded, context-aware responses.
- Collaborated on prompt engineering and evaluation workflows to improve output reliability and reduce hallucination across multiple GenAI product features.

PROJECTS

Full builds for RecruitIQ, ClinicalCrew, AutoSEC, DocMind Hybrid, InterviewIQ, InvestIQ AI under Experience above.

DocuMind Enterprise — Production RAG Knowledge Assistant

LangChain, OpenAI GPT-4o, Pinecone, FastAPI, Docker

- Delivered a production RAG system for internal knowledge search with hybrid retrieval, reranking, and citation-grounded generation; deployed with monitoring/caching for sub-second latency at concurrent load.

OpsFlow AI — Business Process Automation Agent

CrewAI, n8n, GPT-4o, FastAPI

- Scoped and delivered a multi-agent automation system replacing manual data-entry/reporting with autonomous workflows, cutting weekly reporting effort via automated aggregation and report generation.

Bangla-Chittagonian Dialect Neural Machine Translation (BSc Thesis) *QLoRA, Qwen2.5-3B, Gemma-2B-IT, LLaMA-3.2-3B*

- QLoRA fine-tuning of three base LLMs for Bangla-Chittagonian dialect translation — first neural MT system for this dialect pair, supervised by Ferdous Ara.

Interactive Sales Performance Dashboard

Microsoft Excel — Industrial Training, Brain Machine AI

EDUCATION

BSc in Computer Science & Engineering — BGC Trust University Bangladesh

Jun 2022 – Jun 2026

Thesis: Bangla-Chittagonian Dialect NMT via QLoRA Fine-Tuning (Supervisor: Ferdous Ara)